

GRAPHITE-RUST

Xavier Lange

xrlange@tureus.com

xavier.lange@viasat.com

@tureus

ME

- Independent software contractor
- Computer engineering background
 - Working with ViaSat, Inc. on large scale logging
- Dig the outdoors
- Cat owner

MY FAVORITE LANGUAGES, IN NO PARTICULAR ORDER

- OO
 - Ruby
- Functional
 - Erlang
 - Haskell
- Weird
 - Rust



iangreenleaf 2:47 PM

Come on, how many languages can you possibly juggle? This is ridiculous.



xrl 2:47 PM

watch me 😎



iangreenleaf 2:47 PM

I swear @xrl just picks a Language Of The Week



xrl 2:48 PM

I deliver solutions not code 😎 (edited)



iangreenleaf 2:48 PM

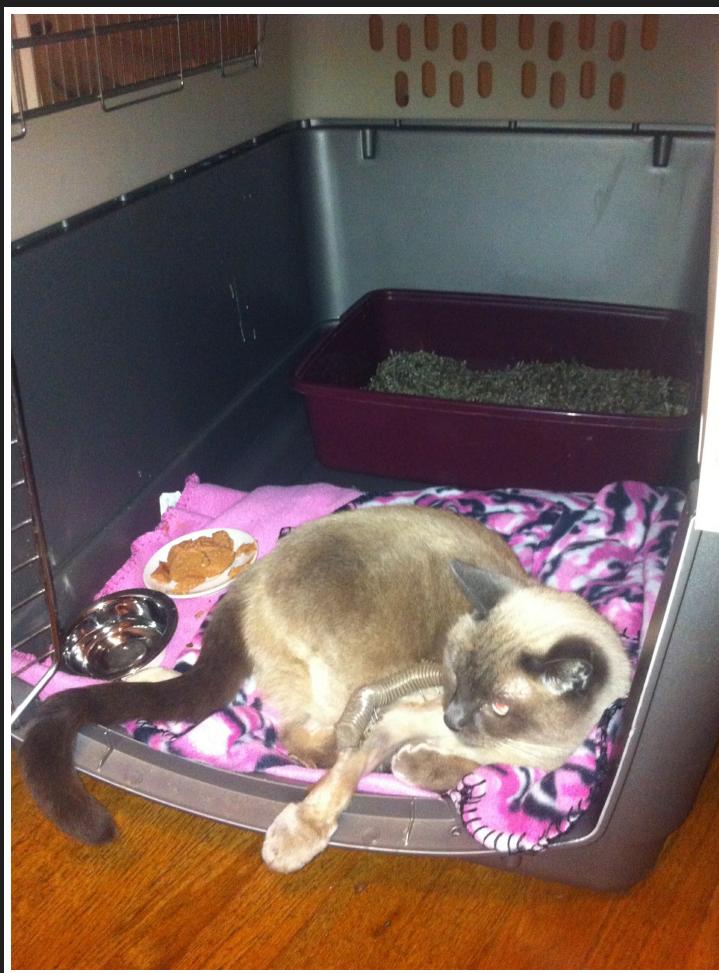
solutions, in whichever language is LOTW this week

MY CAT LUA

Inspiration to stop lounging and go for it



Even if you might get hit by a car



Lua watching standup comedy



GRAPHITE-RUST

- Take a "slow" Python implementation and clone in Rust
- Built to satisfy some needs at ViaSat, Inc.
 - Love graphite
 - Dislike it's opaque performance configuration
 - Troublesome installation
- A dash of NIH
- Supports a private cloud installation and large-ish logging/indexing cluster

TIME SERIES AS A RUST TYPE

Pseudo code

```
fn main() {
    let retention = 60*60*24; // one day
    let mut archive : Vec<(u32,f64)> = Vec::with_capacity(retention);

    let my_series = &mut archive[..];

    my_series[0] = (1438460243,100.0);
}
```

WE'LL COVER

- Graphite concepts (run don't walk)
- Analysis of graphite arch shortcomings (brisk walk)
- The Rust-Graphite implementation (some examples, leisurely strole)
- Analysis of graphite-rust shortcomings (leisurely strole)

LET'S TALK
METRICS

NOT

- Analytics
 - Urchin/Google Analytics
- Business Intelligence
 - Hadoop/Storm
- High-level
 - Web session heatmaps
- Aggregation
 - StatsD

GRAPHITE METRIC CONCEPTS

- Long-lived time series
 - Ex: Measure server A's CPU for 2 years
- Regular update intervals
 - Metric is offered every X seconds (our happy-ish spot: 60s)
- Pre-allocated space on disk
- Basic metric naming
 - one-dimensional
 - ad-hoc

NAMES

Metric names are primitive

```
proglangs.rust.community.active_members  
proglangs.rust.community.lifetime_questions  
proglangs.go.community.active_members  
proglangs.go.community.segfault_questions
```

NAMES

Finding your metrics can get fancy

```
proglangs.*.community.active_members
```

Yields

```
-> proglangs.rust.community.active_members  
-> proglangs.go.community.active_members
```

VALUES

primitive like whoa

(u32 , f64)

VALUES

u32: seconds since epoch. SECONDS

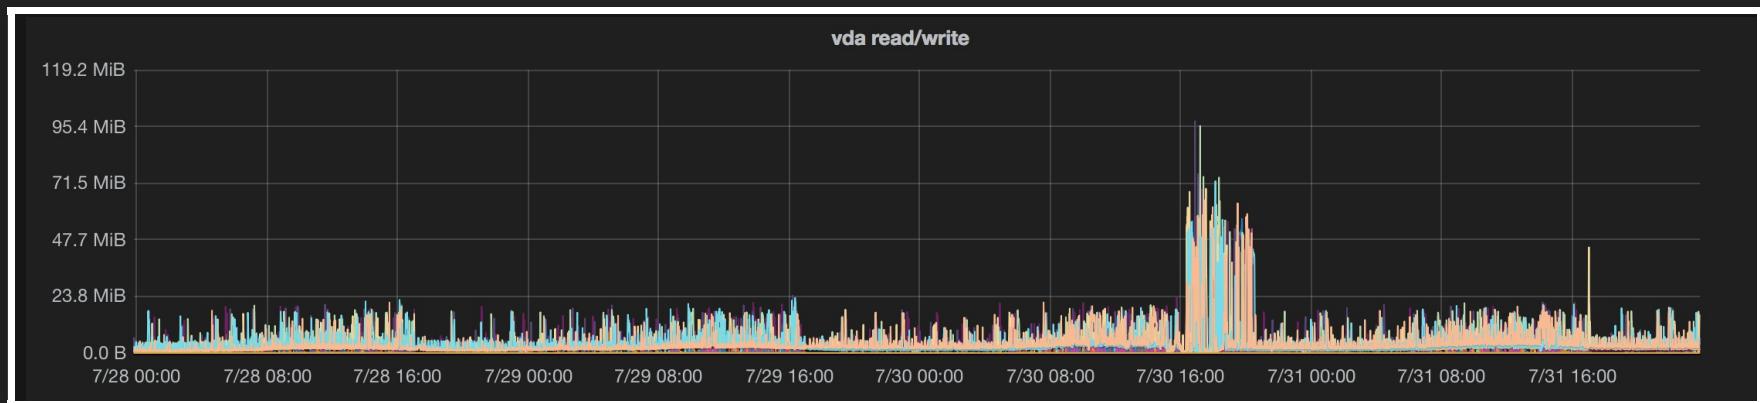
f64: a number, big or small or in between

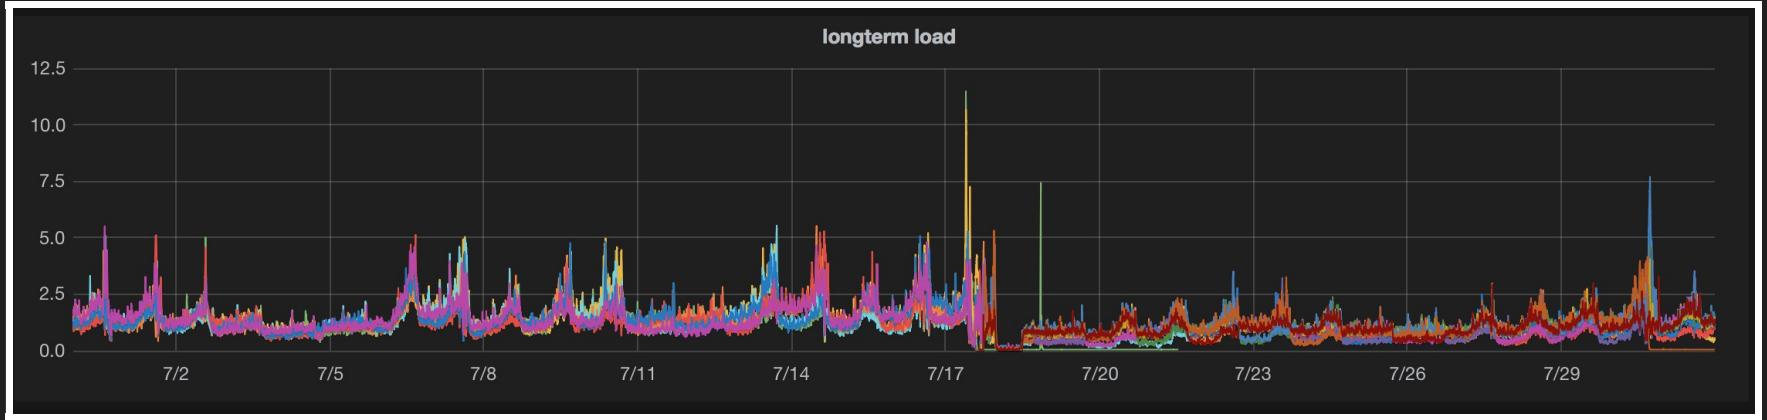
VALUES

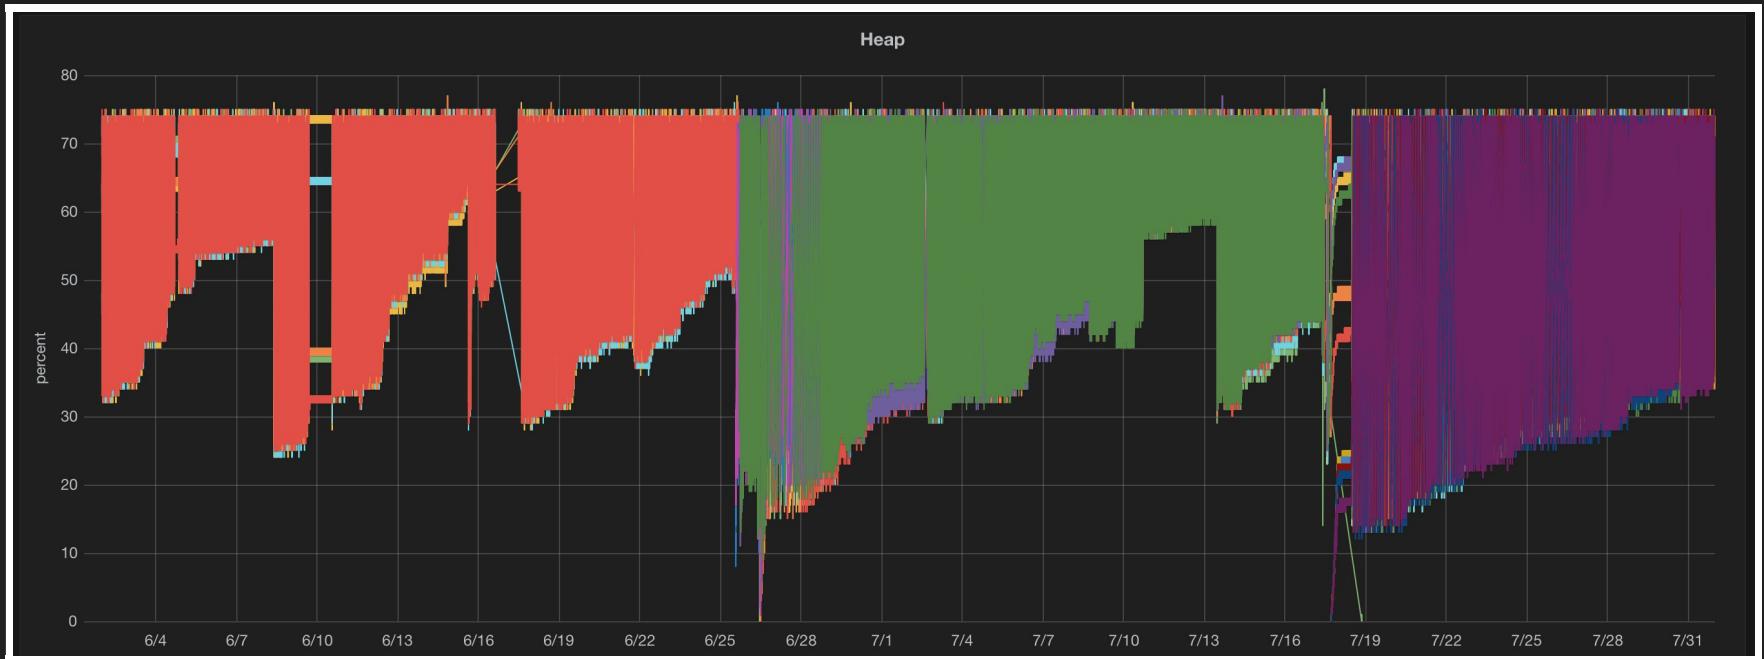
but what can you possibly store with that?

...

enough to be dangerous!







GRAPHITE IN PYTHON

Actually 3 core components, whose names I have adopted

1. Graphite-Web
2. Carbon
3. Whisper

GRAPHTE IN PYTHON INSTALLATION

Tricky, many components

Run `sudo pip install` or learn the `virtualenv`
toolchain

3. GRAPHITE-WEB

A multi-format web service written on top of Django.

- Least interesting to me
- Read-only access to database
- HTML interface for rendering queries
- REST-y interface
 - discovering metrics
 - querying/transform and delivering
 - JSON, CSV, plaintext, pickle, etc
- Tons of features, hidden gems

2. CARBON

A TCP/UDP daemon for recording datapoints

- Moderately interesting to me
- Creates metric storage if necessary
- Writes metrics to disk
- Handles the life-cycle of resources when writing

3. WHISPER

- Fascinating
- Understands bytes on disk
- Writes points
 - Through all archives
- Should be easier than writing a ACID MVCC
 - Like, way easier
 - No real query planner
 - No multi-block joins
 - No transactions

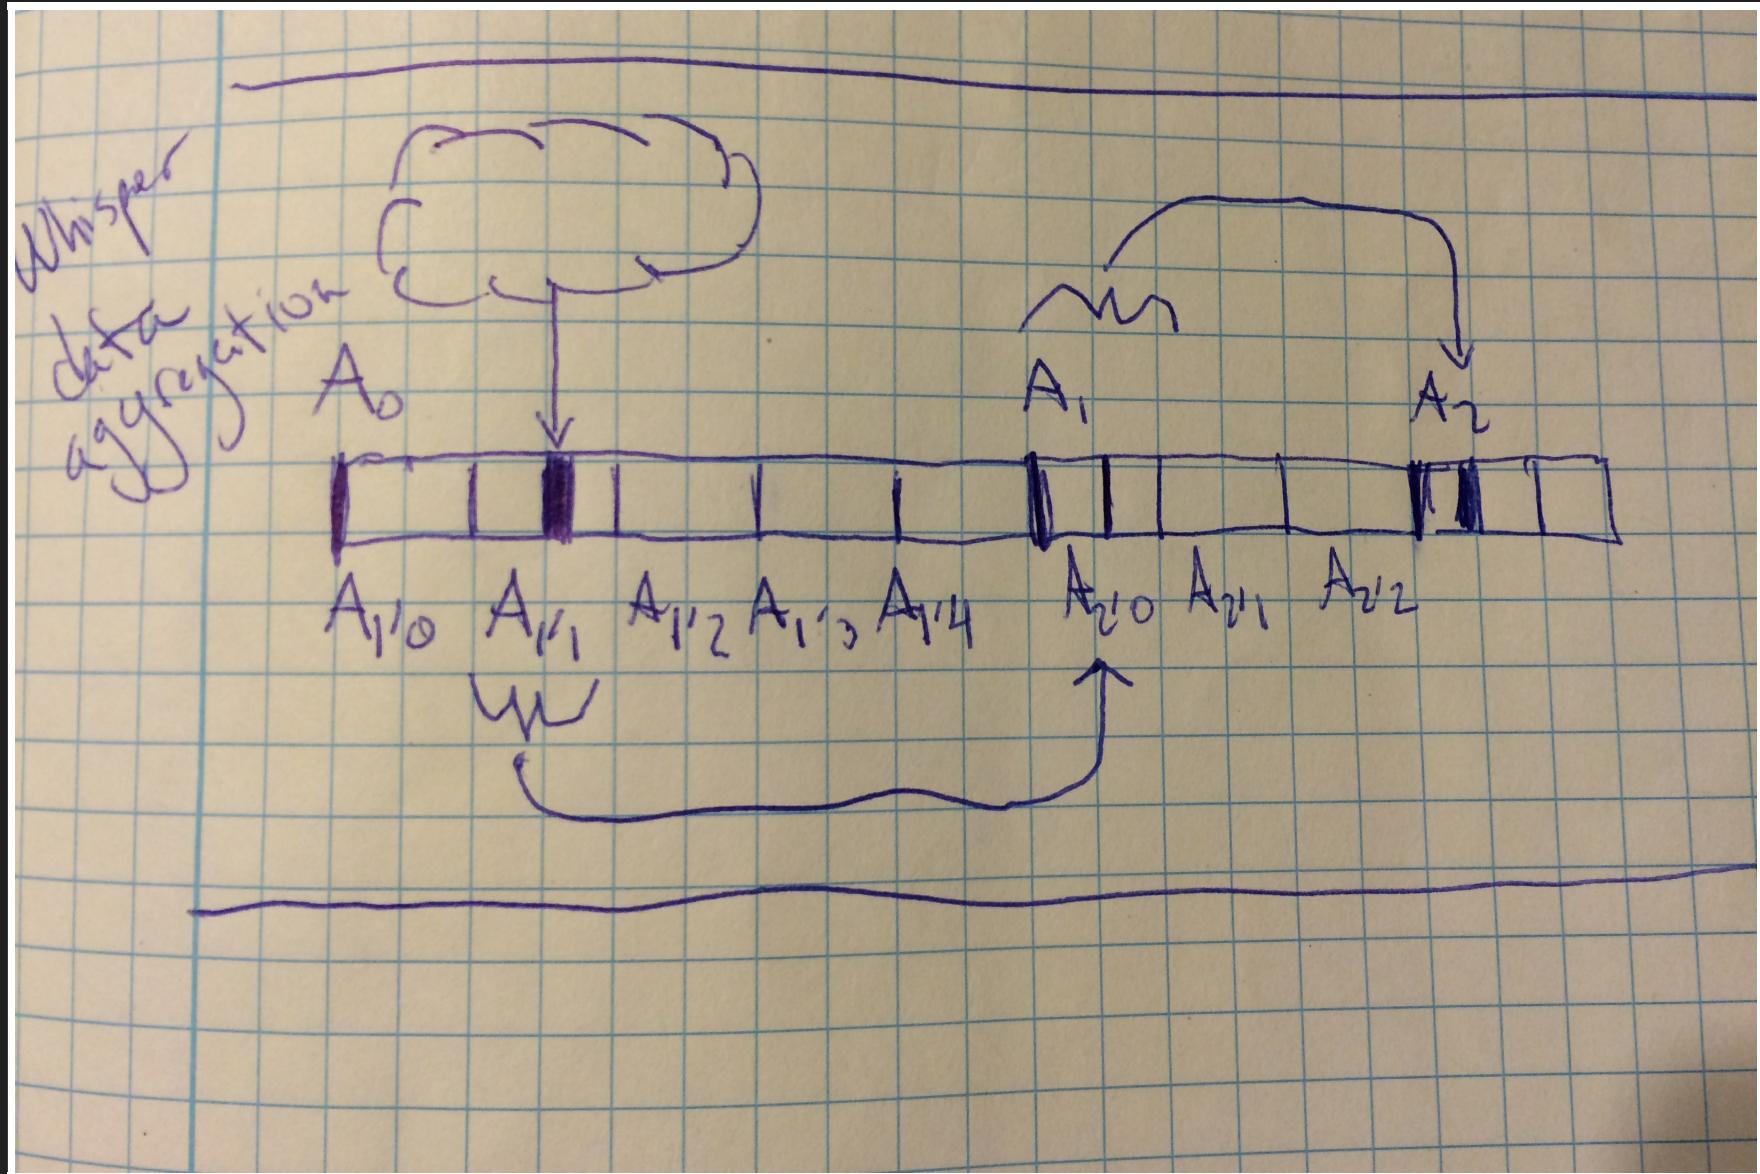
3. WHISPER

... Rust is fast. Maybe that'll make perf automatically faster?

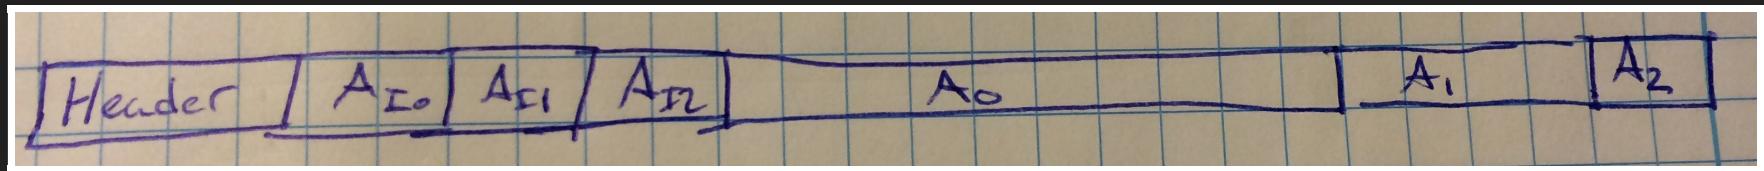
Turned out to be true

DIGGING IN TO WHISPER

VALUE PROPAGATION



INDIVIDUAL ARCHIVE



```
pub trait WhisperFile {
    fn open(&Path) -> Result<Self, Error>;
    fn new(path: &Path, schema: Schema) -> Result<Self, Error>;
    fn write(&mut self, current_time: u64, point: point::Point);
}
```

```
pub struct MutexWhisperFile {
    pub handle: Mutex<File>,
    pub header: Header
}
```

```
// TODO: Don't think we need Copy/Clone.  
// Just added it to make tests easier to write.  
#[derive(PartialEq,Copy,Clone,Debug)]  
pub struct ArchiveInfo {  
    pub offset: SeekFrom,  
    pub seconds_per_point: u64,  
    pub points: u64,  
    pub retention: u64,  
}
```

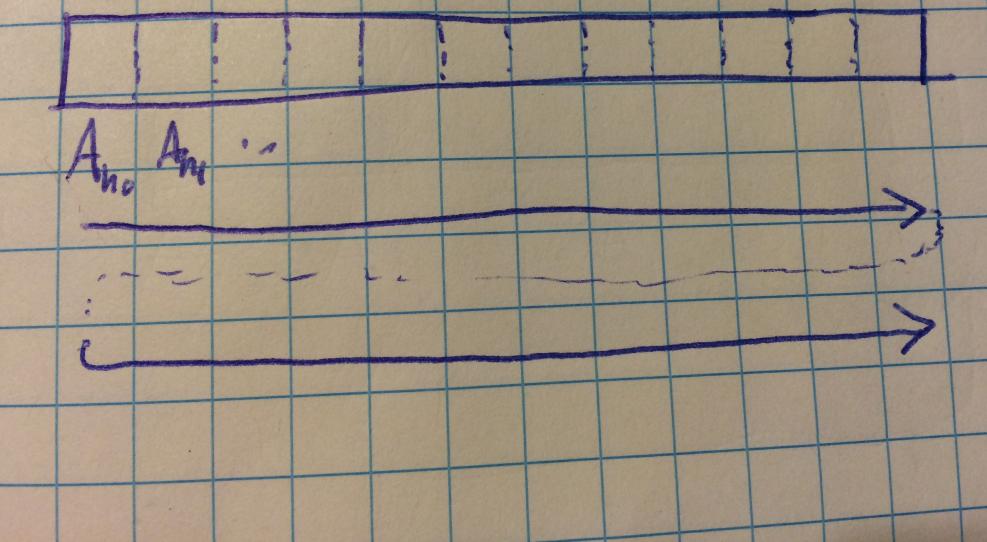
TUPLE STRUCTS

Whisper constructs not present in Python

```
// Index in to an archive, 0..points.len()
#[derive(Debug, PartialEq, PartialOrd)]
pub struct ArchiveIndex(pub u64);

// A normalized timestamp
pub struct BucketName(pub u64);
```

Archive = Circular Buffer



CASE STUDY: WRAP AROUND READ

Used in downsampling

1. Example of "read" amplification
2. I like "pure" code
3. I was figuring out tangled python
4. I wanted to reuse a buffer

CASE STUDY: WRAP AROUND READ FOR DOWNSAMPLING

Type signatures are neat!

```
pub fn read_ops<'a> (h_res_archive: &ArchiveInfo,
                      l_res_archive: &ArchiveInfo,
                      h_res_points: &'a mut [point::Point],
                      point_timestamp: u64,
                      h_res_anchor: BucketName)
-> (
            (ArchiveIndex, &'a mut [point::Point]),
            Option<(ArchiveIndex, &'a mut [point::Point])>
        )
```

CASE STUDY: WRAP AROUND READ

Powered by a favorite: `split_at_mut`

```
// Contiguous read. The easy one.
if h_res_start_index < h_res_end_index {
    ((h_res_start_index, &mut h_res_points[..]), None)
// Wrap-around read
} else {
    let split_index = (hres.points - start_index.0) as usize;
    let (first_buf, second_buf) = points.split_at_mut( split_index
    let zero_index = ArchiveIndex(0);
    ((h_res_start_index, first_buf), Some((zero_index, second_buf))
}
```

IS IT FASTER?

- Yes. Duh.
- 2-3x faster.
- Python: spends more time in userland
- Rust: spends more time in syscalls
- With almost the same naïve behavior as python
- Poor buffer reuse
- Open file, read headers, close it right away

IS IT EASIER TO DISTRIBUTE?

Yes. Duh.

```
$ du -hs target/release/{whisper,carbon,graphite}  
1.1M  target/release/whisper  
1.2M  target/release/carbon  
2.0M  target/release/graphite
```

WHERE DO WE STAND IN
FEATURES TODAY?

CARBON

Uses of the WhisperCache

- Hold files descriptors for longer
- Still need eviction
- Make re-use fast
- Does periodic fsync
- Hides creation/reuse

CARBON'S CACHE

```
pub struct Cache {  
    pub base_path: PathBuf,  
    open_files: HashMap< PathBuf, Box< MutexWhisperFile > >  
}
```

GRAPHITE

Not very usable. Having issues with my libraries.

You can still use the Python implementation.

NEXT STEPS

Keep improving write performance

- mmap data, manual
- get better at dtrace
- carry buffers around

WAS RUST A GOOD CHOICE?

I don't like to advocate rewrites. Especially since I didn't write the first impl.

Rust has great tooling, tests, benchmarks.

I can use mmap.

I can reuse buffers.

I have a compiler yell at me.

So I'm happy.

ANYWAY

Rendering a ton of metrics is my current bottleneck in JS
land

Server side rendering to PNG could be interesting

THANKS

twitter: @tureus

email: xrlange@gmail.com

graphite-rust: <https://github.com/tureus/graphite-rust>